

# BDA Codes

## Steps to execute MapReduce Programs

1. Create one folder on Desktop-"WordCountTutorial"
  - Paste WordCount.java file
  - Create a folder named "Input\_Data" & Create Input.txt file (enter some words)
  - Create a folder name "tutorial\_classes"
2. export HADOOP\_CLASSPATH=\$(hadoop classpath)
3. echo \$HADOOP\_CLASSPATH
4. hadoop fs -mkdir /WordCountTutorial
5. hadoop fs -mkdir /WordCountTutorial/input
6. hadoop fs -put /home/vignan/WordCountTutorial/Input\_Data/Input.txt /WordCountTutorial/input
7. Change the current directory to the tutorial directory  
cd '/home/vignan/WordCountTutorial';  
vignan@vignan-HP-Compaq-8200-Elite-SFF-PC:~/vignan/WordCountTutorial\$
8. Compile the java code  
javac -classpath \${HADOOP\_CLASSPATH} -d /home/vignan/WordCountTutorial/tutorial\_classes /home/vignan/WordCountTutorial/WordCount.java
9. Put the output files in one JAR file  
jar -cvf firstTutorial.jar -C tutorial\_classes/ .
10. Run JAR file  
hadoop jar /home/vignan/WordCountTutorial/firstTutorial.jar WordCount /WordCountTutorial/input /WordCountTutorial/output
11. See the output  
hadoop dfs -cat /WordCountTutorial/output/\*

## Question-1

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import java.io.IOException;

public class StudentGradeClassification {

    // Mapper Class
    public static class GradeMapper extends Mapper<Object, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text grade = new Text();

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
            // Input format: StudentID, Marks
```

```

String[] tokens = value.toString().split(",");
if (tokens.length == 2) {
    try {
        int marks = Integer.parseInt(tokens[1].trim());
        if (marks >= 85) {
            grade.set("A");
        } else if (marks >= 70) {
            grade.set("B");
        } else if (marks >= 50) {
            grade.set("C");
        } else {
            grade.set("D");
        }
        context.write(grade, one);
    } catch (NumberFormatException e) {
        System.err.println("Invalid marks entry: " + tokens[1]);
    }
}
}
}

```

// Reducer Class

```

public static class GradeReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

```

// Driver Class

```

public static void main(String[] args) throws Exception {
    if (args.length != 2) {
        System.err.println("Usage: StudentGradeClassification <input path> <output path>");
        System.exit(-1);
    }
}

```

```

Configuration conf = new Configuration();

```

```

Job job = Job.getInstance(conf, "Student Grade Classification");

```

```

job.setJarByClass(StudentGradeClassification.class);

```

```

job.setMapperClass(GradeMapper.class);

```

```

job.setCombinerClass(GradeReducer.class);

```

```

job.setReducerClass(GradeReducer.class);

```

```

job.setOutputKeyClass(Text.class);

```

```

        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

## Question-2

```

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

public class HighestSalaryEmployee {

    // Mapper Class
    public static class SalaryMapper extends Mapper<Object, Text, Text, Text> {
        private Text salary = new Text();
        private Text employeeDetails = new Text();

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
            // Input format: eid, name, age, salary
            String[] tokens = value.toString().split(",");
            if (tokens.length == 4) {
                String eid = tokens[0].trim();
                String name = tokens[1].trim();
                String age = tokens[2].trim();
                String salaryValue = tokens[3].trim();

                // Emit salary as the key and employee details as the value
                salary.set(salaryValue);
                employeeDetails.set(eid + "," + name + "," + salaryValue);
                context.write(new Text("HighestSalary"), employeeDetails);
            }
        }
    }

    // Reducer Class
    public static class SalaryReducer extends Reducer<Text, Text, Text, Text> {
        private Text result = new Text();

        public void reduce(Text key, Iterable<Text> values, Context context)
            throws IOException, InterruptedException {

```

```

String highestSalaryEmployee = "";
double highestSalary = 0.0;

// Compare salaries to find the highest
for (Text val : values) {
    String[] tokens = val.toString().split(",");
    if (tokens.length == 3) {
        double salary = Double.parseDouble(tokens[2].trim());
        if (salary > highestSalary) {
            highestSalary = salary;
            highestSalaryEmployee = val.toString();
        }
    }
}

result.set(highestSalaryEmployee);
context.write(new Text("Employee with Highest Salary:"), result);
}
}

// Driver Class
public static void main(String[] args) throws Exception {
    if (args.length != 2) {
        System.err.println("Usage: HighestSalaryEmployee <input path> <output path>");
        System.exit(-1);
    }

    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "Find Highest Salary Employee");

    job.setJarByClass(HighestSalaryEmployee.class);
    job.setMapperClass(SalaryMapper.class);
    job.setReducerClass(SalaryReducer.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

## Question-6

```

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;

```

```

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

public class CustomerTransactionCount {

    // Mapper Class
    public static class TransactionMapper extends Mapper<Object, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text customerID = new Text();

        @Override
        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
            // Split the input line by spaces
            String[] fields = value.toString().split("\\s+");

            // Extract Customer ID (assumes it's the first field in the record)
            if (fields.length >= 4) { // Ensure there are enough fields in the record
                customerID.set(fields[0]); // Customer ID is in the first field
                context.write(customerID, one); // Emit <Customer ID, 1> for each transaction
            }
        }
    }

    // Reducer Class
    public static class TransactionReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable totalTransactions = new IntWritable();

        @Override
        public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,
        InterruptedException {
            int sum = 0;

            // Sum the values (which are all 1's)
            for (IntWritable val : values) {
                sum += val.get();
            }

            totalTransactions.set(sum); // Set the total transaction count
            context.write(key, totalTransactions); // Emit <Customer ID, Total Transactions>
        }
    }

    // Driver Code
    public static void main(String[] args) throws Exception {
        // Check if proper arguments are passed
        if (args.length != 2) {
            System.err.println("Usage: CustomerTransactionCount <input path> <output path>");
            System.exit(-1);
        }
    }
}

```

```

// Set up the configuration
Configuration conf = new Configuration();
Job job = Job.getInstance(conf, "Customer Transaction Count");

// Set the jar file for the job
job.setJarByClass(CustomerTransactionCount.class);

// Set Mapper and Reducer classes
job.setMapperClass(TransactionMapper.class);
job.setReducerClass(TransactionReducer.class);

// Set the output key and value types
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

// Set input and output paths
FileInputFormat.addInputPath(job, new Path(args[0])); // Input file path
FileOutputFormat.setOutputPath(job, new Path(args[1])); // Output directory path

// Run the job
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

## Question-9

```

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

public class RestaurantRating {

    // Mapper Class
    public static class RatingMapper extends Mapper<Object, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text ratingCategory = new Text();

        @Override
        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
            // Split the input line by space or tab
            String[] fields = value.toString().split("\\s+");

            // Ensure the input has the rating field
            if (fields.length >= 2) {
                int rating = Integer.parseInt(fields[1]);
            }
        }
    }
}

```

```

// Classify the rating into categories
if (rating == 1) {
    ratingCategory.set("Very Poor");
} else if (rating == 2) {
    ratingCategory.set("Poor");
} else if (rating == 3) {
    ratingCategory.set("Average");
} else if (rating == 4) {
    ratingCategory.set("Good");
} else if (rating == 5) {
    ratingCategory.set("Excellent");
} else {
    return; // If rating is out of range, we ignore this record
}

// Emit the rating category and a count of 1 for each rating occurrence
context.write(ratingCategory, one);
}
}
}

```

// Reducer Class

```

public static class RatingReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable totalRatings = new IntWritable();

    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,
    InterruptedException {
        int sum = 0;

        // Sum the values (each value is 1, so we are counting occurrences)
        for (IntWritable val : values) {
            sum += val.get();
        }

        totalRatings.set(sum); // Set the total count for the rating category
        context.write(key, totalRatings); // Emit the category and its total count
    }
}

```

// Driver Code

```

public static void main(String[] args) throws Exception {
    // Check if proper arguments are passed
    if (args.length != 2) {
        System.err.println("Usage: RestaurantRating <input path> <output path>");
        System.exit(-1);
    }
}

```

// Set up the configuration

```

Configuration conf = new Configuration();
Job job = Job.getInstance(conf, "Restaurant Rating Classification");

```

```

// Set the jar file for the job
job.setJarByClass(RestaurantRating.class);

// Set Mapper and Reducer classes
job.setMapperClass(RatingMapper.class);
job.setReducerClass(RatingReducer.class);

// Set the output key and value types
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

// Set input and output paths
FileInputFormat.addInputPath(job, new Path(args[0])); // Input file path
FileOutputFormat.setOutputPath(job, new Path(args[1])); // Output directory path

// Run the job
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

## Steps to enter hive shell:

```

start-all.sh
hive

```

## Question-3

Tables creation:

- create table if not exists Doctors (DoctorName string, Specialization string, ExperienceYears int, ContactNumber string) row format delimited fields terminated by ',' stored as a textfile;
- create table if not exists Patients (PatientOpNumber string, Name string, Age int, Gender string, Contact string) row format delimited fields terminated by ',' stored as a textfile;

Data Insertion: Insert some sample data in local textfiles with comma separated values.

- load data local inpath 'path\_to\_doctors.txt' into table Doctors;
- load data local inpath 'path\_to\_patients.txt' into table Patients;

Verify data :

- select \* from Doctors;
- select \* from Patients;

## Question-4

Table creation:

- create table if not exists Instagram(PostId int, Likes int, Comments int) row format delimited fields terminated by ',' stored as textfile;

Data insertion: Insert some data in a local txtfile with comma separated values.

- Load data local inpath 'path\_to\_instagram.txt' into table Instagram;

Queries:

- i. select count(PostId) as TotalPosts from Instagram;
- ii. select avg(Likes) as AverageLikes, avg(Comments) as AverageComments from Instagram;

## Question-11

Table creation

- create table if not exists Books(BookId int, Title string, Author string, Genre string, Rating float) row format delimited fields terminated by ',' stored as textfile;
- describe Books;

Data insertion: Insert some data in a local txtfile with comma separated values.



- Load data local inpath'path\_to\_books.txt' into table Books;

Queries:

- select \* from Books where Rating > 4.0;
- select \* from Books where Genre != 'Fiction';

## Steps to enter grunt shell:

start-all.sh

pig

Question-5

## Question-13:

# Step 1: Install PySpark and set up environment

```
!apt-get install openjdk-17-jdk-headless -qq > /dev/null
```

```
!wget -q https://dlcdn.apache.org/spark/spark-3.5.3/spark-3.5.3-bin-hadoop3.tgz
```

```
!tar xf spark-3.5.3-bin-hadoop3.tgz
```

```
!pip install -q findspark
```

```
import os
```

```
import findspark
```

```
findspark.init()
```

# Step 2: Set environment variables for Spark and Hadoop

```
os.environ["SPARK_HOME"] = "/content/spark-3.5.3-bin-hadoop3" # Corrected path for Spark 3.5.3
```

```
os.environ["JAVA_HOME"] = "/usr/lib/jvm/java-17-openjdk-amd64" # Corrected to Java 17
```

# Step 3: Initialize SparkContext

```
from pyspark.sql import SparkSession
```

# Ensure SPARK\_HOME is set correctly and without '/bin'

```
spark = SparkSession.builder.master("local[*]").appName("TCS Selection").getOrCreate()
```

```
sc = spark.sparkContext
```

# Step 4: Sample data for selected students (student\_id, name)

```
data = [
```

```
    (1, "John"),
```

```
    (2, "Alice"),
```

```
    (3, "Bob"),
```

```
    (4, "Charlie"),
```

```
    (5, "David")
```

```
]
```

# Create RDD

```
students_rdd = sc.parallelize(data)
```

# Step 5: Count total number of students selected for TCS

```
total_students = students_rdd.count()
```

```
print(f"Total number of students selected for TCS: {total_students}")
```

# Step 6: Sort the selected students by name in ascending order

```
sorted_students_rdd = students_rdd.sortBy(lambda x: x[1]) # Sorting by name (index 1)
```

```
sorted_students = sorted_students_rdd.collect()
```

```
print("Sorted students by name (ascending):")
```

```
for student in sorted_students:
```

```
    print(student)
```